

How to read each story (template)

- **Context:** who, what they do, where the manual work is today
- **Pain:** short list of issues worth money
- **To-be workflow (user journey):** human steps + what the Pipeline Manager does
- **Pipeline (blocks):** Sources → Ingest & Parse → Understand → Verify → Orchestrate → Outputs
- **Governance:** residency, PII handling, audit trail
- **KPIs / targets (pilot-year):** what we measure (align with your KPI slide)
- **Deploy & tier fit:** SaaS vs On-prem, Basic / Pro / Enterprise
- **Image brief:** one-liner for a 16:9 diagram in your “HOW IT WORKS” style

A) Personal / Solo

Story A1 — Independent Field Inspector (images + PDFs → verified reports)

Context

A single-person contractor does factory site inspections for OEMs. They take photos/video on site and write reports later.

Pain

- Nights lost compiling photos, notes, and standards references
- Missed non-conformities; inconsistent wording; slow report turnaround

To-be workflow (user journey)

1. Inspector uploads **photos/video** from phone; drags the **customer checklist PDF**.
2. **Pipeline Manager** proposes a prebuilt “Inspection Report” pipeline.
3. Vision agent auto-labels findings; doc agent drafts the report; verify step flags any missing photos for a checklist item.
4. Inspector reviews, adds comments, **e-signs**, shares to client portal.

Pipeline (blocks)

- **Sources:** Images, Video, Checklist PDFs, Past reports

- **Ingest & Parse:** OCR, captions, embeddings → **Buckets**
- **Understand:** Vision (defect tagging), Embedding + K/V extraction
- **Verify:** Checklist coverage, photo-to-item linkage, date/location EXIF match
- **Orchestrate:** Route to “Report Gen” agent → export PDF + JSON
- **Outputs:** Signed report, defect gallery, CSV of findings

Governance

Local device encryption; no PII; keep 180-day retention.

KPIs / targets (pilot-year)

- Turnaround time: **48–72h → same day**
- Missed items: **–60%**; Client acceptance first pass: **≥90%**

Deploy & tier fit

SaaS, **Basic**. Usage metering: media minutes + docs.

Image brief

“Phone **Images/Video** + **Checklist PDF** → **Ingestion** → **Vision tag** → **Verify checklist** → **Generate report** → 🤖 client.”

B) Small Business (≈10–50 FTE)

Story B1 — HR Consultancy (CVs & compliance packs → shortlist + audit trail)

Context

Small HR consultancy screens 500+ CVs/month for industrial clients; each hire requires a compliance document pack.

Pain

- Manual CV parsing; inconsistent screening notes
- Compliance packs assembled by hand; no lineage

To-be workflow

1. Drag & drop **CVs (PDF/Doc)** + **job spec**.
2. Agent extracts skills, history, red flags; ranks candidates.
3. Compliance pack generator collects ID/proof docs, validates against checklist.
4. Client receives a **shortlist** + audit report.

Pipeline (blocks)

- **Sources:** CV PDFs, Job description, Identity docs, ERP/BPM job reqs
- **Ingest:** OCR/Parsers → **Buckets**
- **Understand:** K/V + entity extraction; embeddings; similarity to job spec
- **Verify:** Required docs present; dates/IDs match; duplicate check
- **Orchestrate:** Ranker agent → Pack builder agent → E-signature
- **Outputs:** Ranked shortlist, compliance pack PDF, lineage graph

Governance

PII masked at rest; DLP for ID numbers; access by client project.

KPIs / targets

- Time per req: **6h → 1h**
- Pack completeness on first try: **≥95%**
- Audit finding rate: **0**

Deploy & tier fit

SaaS, **Pro** (needs 3–5 agents, higher model quota).

Image brief

“**CVs + Job Spec** → Parse → **Rank** → **Verify** (docs checklist) → **Pack** → **Client**.”

Story B2 — Precision Machine Shop (work orders + images → traceable QA)

Context

Small shop makes custom metal parts; customers require **traceability** (work order ↔ photo evidence).

Pain

- Photos not linked to work order/step; QA PDF later assembled by hand
- Rework due to missing proof

To-be workflow

Operator scans QR on work order; takes photos per step; pipeline binds media → step → requirement; generates QA dossier.

Pipeline

- **Sources:** ERP work orders, Step checklists, Images/Video
- **Ingest:** OCR/metadata; **Buckets**
- **Understand:** Vision tags (finish, dimensions screenshot), entity linking to order/step
- **Verify:** Step coverage, timestamp within operation window
- **Orchestrate:** QA dossier generator + dashboard
- **Outputs:** QA dossier PDF, step lineage graph, exceptions queue

Governance

On-prem bucket mirror optional; 3-year retention.

KPIs

- Late dossiers: **-80%**
- Chargeback reductions: **-50%**
- First-pass yield: **+5-10%**

Deploy & tier fit

On-prem (shop floor), **Pro**.

Image brief

“ERP WO → Photos per step → Link → Verify coverage → QA PDF.”

C) Mid-Market (≈100–1,000 FTE)

Story C1 — Pharmaceutical QA (batch records + lab images → compliance pack)

Context

GxP manufacturer assembling **Batch Manufacturing Records (BMRs)** with lab images and equipment logs.

Pain

- Manual review; deviations missed; audit scramble

To-be workflow

1. Connect LIMS, MES, equipment logs, scanned BMR pages, microscope images.
2. Pipeline extracts parameters, links to spec limits, flags deviations.
3. Auto-builds **eCTD-ready** compliance pack with lineage.

Pipeline

- **Sources:** LIMS, MES, PDFs, Images, Sensors/APIs
- **Ingest:** OCR/parsers; **Buckets**
- **Understand:** K/V, entities (lot, operator), embeddings; rule engine for specs
- **Verify:** Cross-validate lot IDs, time windows, spec conformance
- **Orchestrate:** Deviation routing → CAPA draft → Pack generator
- **Outputs:** Deviation log, CAPA draft, compliance pack, audit trail

Governance

On-prem/private cloud; role-based access; immutable audit log.

KPIs

- Review time per batch: **-60-70%**
- Deviation detection sensitivity: **+30%**
- Audit observations: **-50%**

Deploy & tier fit

On-prem, Enterprise.

Image brief

“LIMS/MES + PDFs + Images + Sensors → Extract → **Verify specs** → **Route CAPA** → **Compliance Pack.**”

Story C2 — 3PL Logistics (POD images + EDI docs → dispute automation)

Context

A regional 3PL handles Proof-of-Delivery (POD) photos, signed slips, and **EDI 214/210**.

Pain

- Chargebacks from “missing proof”; manual reconciliation across EDI + images

To-be workflow

1. Ingest POD photos, signed slips, EDI feeds.
2. Cross-link shipment → stop → document; verify SLA windows.
3. Auto-reply pack for disputes; dashboard of at-risk shipments.

Pipeline

- **Sources:** EDI, Photos, PDFs, TMS API
- **Ingest:** OCR/embeddings; **Buckets**
- **Understand:** Entity linking (BOL, PRO, stop), vision “signature present”
- **Verify:** SLA window, geo/time consistency, completeness
- **Orchestrate:** Dispute pack generator → email/API submit
- **Outputs:** Pack PDFs, exceptions queue, ops dashboard

Governance

Data separation per shipper; 1-year retention.

KPIs

- Dispute cycle time: **-70%**
- Chargeback recovery: **+30-50%**
- Manual touches per shipment: **-60%**

Deploy & tier fit

SaaS or virtual private cloud, **Pro**.

Image brief

“**EDI + POD photo + Signed slip** → Link → **Verify SLA** → **Auto dispute pack** → \$ recovered.”

D) Enterprise (1,000+ FTE)

Story D1 — Energy Utility (asset inspections video + SCADA → risk scoring)

Context

Utility must inspect transmission assets (drones, trucks) and reconcile with **SCADA** alarms and maintenance history.

Pain

- Petabytes of video; weak linkage to asset & alarm; slow triage

To-be workflow

1. Connect drone video, images, SCADA events, EAM/CMMS work orders.
2. Vision models detect defects (corrosion, vegetation); link to asset registry.
3. Risk model scores assets; Pipeline Manager routes **work orders** automatically.

Pipeline

- **Sources:** Video/Images, SCADA, EAM/CMMS, GIS
- **Ingest:** Stream → captions + keyframes; **Buckets**
- **Understand:** Vision (defect classes), embeddings; graph link to asset & history
- **Verify:** GPS/asset match; alarm correlation within window
- **Orchestrate:** Risk scoring agent → WO creation via API → notify crews
- **Outputs:** Ranked risk dashboard, WO list, evidence clips

Governance

Strict residency; network-isolated **on-prem** cluster; auditing.

KPIs

- Triage time: **days → hours**
- Unplanned outages: **–10–20%**
- Truck rolls avoided: **–15%**

Deploy & tier fit

On-prem, Enterprise (custom models, SLOs).

Image brief

“**Drone Video + SCADA + EAM** → Detect → **Link to asset graph** → **Risk score** → **Auto work order.**”

Story D2 — Global Manufacturer (AP/Procurement docs + ERP → touchless matching)

Context

Multi-ERP landscape; invoices, POs, GRNs across regions.

Pain

- Low straight-through match; duplicate vendors; compliance leakage

To-be workflow

1. Pipelines ingest invoices (PDF/EDI), POs, GRNs, vendor master.
2. Extraction + fuzzy matching; anomaly rules; human review only for exceptions.
3. Continuous vendor deduping; compliance checks (tax IDs, sanctions).

Pipeline

- **Sources:** ERP(s), Invoices PDF/EDI, Vendor master
- **Ingest:** Parsers/OCR → **Buckets**
- **Understand:** Line-item extraction, entity resolution, embeddings
- **Verify:** 3-way match (PO/GRN/Invoice), tax rules, duplicate detection
- **Orchestrate:** Exceptions queue, auto-post to ERP, monthly compliance report
- **Outputs:** Posted vouchers, exceptions, audit & savings dashboard

Governance

Privately peered cloud or on-prem; SOX audit trail; PII masking.

KPIs

- Straight-through processing: **+40–60%**
- Duplicates & overpayments: **–70%**
- Cycle time: **–50%**

Deploy & tier fit

VPC or on-prem, **Enterprise**.

Image brief

“Invoices + PO + GRN + Vendor master → Extract → Verify 3-way → Exceptions → Auto-post.”